

JAXMg Basics

Introduction to sharding and multi-GPU linear solvers

1 Introduction

JAXMg wraps NVIDIA’s **cuSolverMg** (multi-GPU solver library) for JAX, enabling distributed linear algebra operations across multiple GPUs. This tutorial covers the essential concepts for using JAXMg effectively.

Function	Solves	Complexity
potrs	$Ax = b$ via Cholesky (A symmetric positive definite)	$O(N^3/3)$
potri	A^{-1} via Cholesky	$O(N^3)$
syevd	Eigendecomposition of symmetric matrix	$O(N^3)$

Table 1: JAXMg distributed linear algebra functions

2 Core Concepts: JAX Sharding

2.1 The Mesh

A **mesh** is a logical n -dimensional grid of devices with named axes. Think of it as organizing your GPUs into a coordinate system that JAX uses to distribute data and computation.

```
1 import jax
2 from jax.sharding import Mesh, NamedSharding, PartitionSpec as P
3
4 # 1D Mesh: 4 GPUs in a line, axis named "x"
5 devices = jax.devices("gpu") # [gpu:0, gpu:1, gpu:2, gpu:3]
6 mesh_1d = jax.make_mesh((4,), ("x",))
7
8 # 2D Mesh: 4 GPUs as 2x2 grid, axes "S" and "T"
9 import numpy as np
10 mesh_2d = Mesh(np.array(devices).reshape(2, 2), ("S", "T"))
```




Figure 1: Device mesh configurations. Left: 1D mesh with 4 GPUs along axis “x”. Right: 2D mesh with 4 GPUs arranged in a 2×2 grid with axes “S” (rows) and “T” (columns).

2.2 PartitionSpec: How to Shard

`PartitionSpec` (aliased as `P`) tells JAX which mesh axis to shard each array dimension along. `None` means that dimension is replicated (not sharded).

PartitionSpec	Meaning (for 2D array of shape (N, M))
<code>P("x", None)</code>	Shard dim 0 (rows) across "x", replicate dim 1 (columns)
<code>P(None, "x")</code>	Replicate dim 0 (rows), shard dim 1 (columns) across "x"
<code>P("S", "T")</code>	Shard rows across "S", columns across "T"
<code>P(None, None)</code>	Fully replicate — each device has a complete copy
<code>P(("S", "T"), None)</code>	Flatten S×T mesh into 1D, shard rows across all devices

Table 2: `PartitionSpec` examples for a 2D array

2.3 Placing Data on Devices


```

1 import jax.numpy as jnp
2
3 N = 1000
4 A = jnp.eye(N, dtype=jnp.float64)    # Matrix: 1000 × 1000
5 b = jnp.ones((N, 1), dtype=jnp.float64)  # Vector: 1000 × 1
6
7 # Create mesh and shard data
8 mesh = jax.make_mesh((4,), ("x",))
9 A_sharded = jax.device_put(A, NamedSharding(mesh, P("x", None)))  # Row-sharded
10 b_replicated = jax.device_put(b, NamedSharding(mesh, P(None, None))) # Replicated

```

Matrix A (1000×1000)

Sharding: `P("x", None)`



Vector b (1000×1)

Sharding: `P(None, None)`




Figure 2: Data distribution across 4 GPUs. Left: Matrix A is row-sharded — each GPU owns 250 rows but all 1000 columns. Right: Vector b is replicated — each GPU has a complete copy.

3 Using `potrs`: Distributed Cholesky Solve

3.1 Basic Usage

```

1 from jaxmg import potrs
2
3 # Problem: Solve Ax = b where A is symmetric positive definite

```

```

4 N = 1000
5 A = jnp.eye(N, dtype=jnp.float64) * 2 # Diagonal matrix
6 b = jnp.ones((N, 1), dtype=jnp.float64)
7
8 # Setup mesh and sharding
9 mesh = jax.make_mesh((4,), ("x",))
10 A = jax.device_put(A, NamedSharding(mesh, P("x", None)))
11 b = jax.device_put(b, NamedSharding(mesh, P(None, None)))
12
13 # Solve!
14 x = potrs(A, b, T_A=128, mesh=mesh, in_specs=(P("x", None), P(None, None)))
15 # x ≈ [0.5, 0.5, 0.5, ...]

```

3.2 Key Parameters

Parameter	Description
T_A	Tile size — block width for distributed algorithm. Larger = faster but more memory. Recommended: 128–2048. Auto-pads if shard size not divisible.
mesh	JAX Mesh object defining device topology
in_specs	Tuple of PartitionSpecs: (P("axis", None), P(None, None)) — matrix row-sharded, RHS replicated
return_status	If True, returns (solution, status_code) tuple
pad	If True (default), auto-pad shards to be divisible by T_A

Table 3: `potrs` function parameters

3.3 Blocked Cholesky and Cyclic Column Sharding

cuSolverMg uses a right-looking blocked Cholesky factorization. With tile size T_A , the matrix is partitioned into $T_A \times T_A$ tiles. The k -th block column (the panel) consists of tiles $A_{(i, k)}$. Each iteration:

- Factor the diagonal block: $A_{kk} \rightarrow L_{kk}L_{kk}^\dagger$ (Cholesky on the current tile).
- Panel solve: solve $L_{ik}L_{kk}^\dagger = A_{ik}$ for L_{ik} (TRSM on the column panel).
- Trailing update: $A_{ij} = A_{ij} - L_{ik}L_{jk}^\dagger$ (GEMM/SYRK on the remaining tiles).

The GPU that owns block column k performs the diagonal factorization and panel solve; then all GPUs apply trailing updates to their local tiles.

JAXMg balances this with a 1D block-cyclic column layout: columns are assigned to devices in round-robin tiles of width T_A instead of contiguous chunks. This spreads panel ownership, reduces idle time, and avoids a single GPU owning all early columns. Without cyclic sharding, early panels land on one GPU, so others wait for panel data before updating. Internally, JAXMg permutes column tiles into cyclic order using peer-to-peer copies and small staging buffers; the redistribution is deterministic and hidden from the JAX API.

3.4 Internal Workflow




Figure 3: Distributed Cholesky solve workflow. The matrix is distributed across GPUs; cuSolverMg handles the distributed factorization and triangular solves with automatic inter-GPU communication.

3.5 Tile Size Selection

The tile size T_A controls the block granularity of the distributed algorithm:

```

1 N = 10000
2 ndev = 4
3 shard_size = N // ndev # 2500 rows per GPU
4
5 # Good choices: T_A divides shard_size evenly (no padding needed)
6 T_A = 250 # 2500 / 250 = 10 tiles ✓
7 T_A = 500 # 2500 / 500 = 5 tiles ✓
8 T_A = 2048 # 2500 / 2048 = 1.22 → padding added automatically △

```



Tips for choosing T_A :

- Use $T_A \geq 128$ for good performance (small tiles = slow)
- Larger $T_A \rightarrow$ fewer communication rounds, but more memory per tile
- Best: choose T_A that divides $N // n_{dev}$ evenly to avoid padding overhead
- Common good choices: 256, 512, 1024, 2048

4 The `shard_map` Pattern

4.1 What is `shard_map`?

`jax.shard_map` lets you write code that runs **per-shard** on each device, with explicit collective operations for cross-device communication:


```
1  from functools import partial
2
3  @partial(jax.shard_map,
4      mesh=mesh,
5      in_specs=P("x", None),    # Input: row-sharded
6      out_specs=P(None))      # Output: replicated
7  def my_distributed_fn(local_A):
8      # This code runs on EACH GPU independently
9      # local_A shape: (N/ndev, M) – just this GPU's shard
10
11     my_idx = jax.lax.axis_index("x") # Which GPU am I? (0, 1, 2, or 3)
12
13     # Collective: sum across all GPUs
14     global_sum = jax.lax.psum(local_sum, axis_name="x")
15
16     # Collective: gather all shards
17     full_A = jax.lax.all_gather(local_A, axis_name="x", axis=0, tiled=True)
18
19     return result
```

4.2 Common Collective Operations

Operation	Description	Output Shape
<code>psum(x, "axis")</code>	Sum values across all devices on axis	Same shape, replicated
<code>pmean(x, "axis")</code>	Mean across devices	Same shape, replicated
<code>all_gather(x, "axis")</code>	Concatenate shards from all devices	Larger (combined)
<code>axis_index("axis")</code>	Get this device's index on axis	Scalar: 0 to ndev-1
<code>pbroadcast(x, "axis")</code>	Broadcast from source device	Same on all devices

Table 4: JAX collective operations for `shard_map`

psum: Sum across devices



all_gather: Collect all shards




Figure 4: Collective operations illustrated. Left: `psum` sums values and replicates result. Right: `all_gather` collects all shards into complete array on each device.

4.3 Using `potrs_shardmap_ctx`

For advanced use inside your own `shard_map` (when you need custom logic before/after the solve):

```
1 from jaxmg import potrs_shardmap_ctx
2
3 def custom_solve(local_A, b):
4     """Runs inside shard_map context - local_A is just this GPU's shard"""
5     local_A = local_A * some_scaling_factor # Custom pre-processing
6     x, status = potrs_shardmap_ctx(local_A, b, T_A=256) # No extra shard_map
7     return x, status
8
9 result = jax.shard_map(
10     partial(custom_solve, b=b_replicated),
11     mesh=mesh,
12     in_specs=P("x", None),
13     out_specs=(P(None, None), P(None)),
14     check_vma=False, # Required for FFI calls
15 )(A_sharded)
```

5 Application: MinSR with JAXMg

5.1 The MinSR Algorithm

In Variational Monte Carlo, **minSR** (minimum-step Stochastic Reconfiguration) computes parameter updates efficiently when $N_s \ll N_p$ (fewer samples than parameters):

$$\delta\theta = \tau \cdot X^\dagger (XX^\dagger + \lambda I)^{-1} \cdot E^{\text{loc}}$$

$X \in \mathbb{R}^{N_s \times N_p}$ Centered Jacobian (samples \times parameters)

$T = XX^\dagger \in \mathbb{R}^{N_s \times N_s}$ **Gram matrix** — much smaller than $N_p \times N_p$!

λ Regularization (diagonal shift, typically 10^{-4} to 10^{-2})

E^{loc} Centered local energy vector

5.2 Multi-Node Mesh Setup

For clusters with multiple nodes, use a 2D mesh where one axis is within-node and another is across-nodes:

```

1 def get_device_grid():
2     """Organize devices: rows = local GPUs, columns = nodes"""
3     by_proc = {}
4     for d in jax.devices():
5         by_proc.setdefault(d.process_index, []).append(d)
6     hosts = sorted(by_proc)
7     return np.array([[by_proc[h][x] for h in hosts]
8                     for x in range(jax.local_device_count())]).T
9
10 def create_2d_mesh():
11     dev_grid = get_device_grid() # shape: (n_nodes, n_gpus_per_node)
12     return Mesh(dev_grid, ("S", "T"))
13 # "S" = across nodes, "T" = within node

```

Multi-Node Setup: 2 Nodes \times 4 GPUs = 8 Total




Figure 5: 2D mesh for multi-node cluster. Axis “S” spans across nodes (inter-node communication via network). Axis “T” spans GPUs within each node (fast NVLink/PCIe). Total 8 devices in a (2, 4) grid.

5.3 MinSR Data Flow

MinSR Computation Pipeline




Figure 6: MinSR data flow. The Jacobian X is sharded by samples. The Gram matrix $T=XX^T$ is computed distributedly. JAXMg's **potrs** solves $(T+\lambda I)v = E^{\text{loc}}$. Final update $\delta\vartheta = X^T v$ uses distributed matrix-vector product with **psum** reduction.

5.4 Key Code Snippets

Sharding the Jacobian:

```
1 0_LT = jax.lax.with_sharding_constraint(
2     0_LT, NamedSharding(mesh_2d, P(("S", "T"), None))
3 ) # Samples sharded across all 8 GPUs, params replicated
```

Distributed Gram matrix (memory-efficient streaming):

```
1 def streamed_gram(0_L, chunk_size):
2     """Compute  $T = X @ X^T$  without materializing full intermediate"""
3     def step(acc, chunk):
4         full_chunk = jax.lax.all_gather(chunk, ('S', 'T'), tiled=True)
5         return acc + chunk @ full_chunk.T, None
6     return jax.lax.scan(step, zeros, 0_L_chunked)[0]
```

Distributed Cholesky solve:

```
1 v = potrs(T_reg, E_loc[:, None], T_A=2048, mesh=mesh_2d,
2           in_specs=(P("T", None), P(None, None)))
```

Final update with reduction:

```
1 def update(0_LT, v):
2     return jax.lax.psum(v @ 0_LT, axis_name=("S", "T"))
3 delta_theta = jax.shard_map(update, mesh_2d,
4     in_specs=(P(("S", "T"), None), P(("S", "T"))), out_specs=P(None))(0_LT, v)
```

6 Complete Example

```
1 import jax
2 import jax.numpy as jnp
3 from jax.sharding import Mesh, NamedSharding, PartitionSpec as P
4 from jaxmg import potrs
5
6 # Setup
7 jax.config.update("jax_enable_x64", True)
8 devices = jax.devices("gpu")
9 ndev = len(devices)
10 mesh = jax.make_mesh((ndev,), ("x",))
11
12 # Create Gram matrix (like in minSR: T = X @ X.T)
13 N_samples, N_params = 4000, 100000
14 key = jax.random.PRNGKey(0)
15 X = jax.random.normal(key, (N_samples, N_params)) / jnp.sqrt(N_samples)
16 T = X @ X.T # (4000, 4000) — much smaller than (100000, 100000)!
17
18 # Add regularization
19 T_reg = T + 1e-4 * jnp.eye(N_samples)
20
21 # RHS (centered local energies)
22 b = jax.random.normal(key, (N_samples, 1))
23
24 # Shard for distributed solve
25 T_sharded = jax.device_put(T_reg, NamedSharding(mesh, P("x", None)))
26 b_replicated = jax.device_put(b, NamedSharding(mesh, P(None, None)))
27
28 # Distributed Cholesky solve: (T + λI) v = b
29 v = potrs(T_sharded, b_replicated, T_A=256, mesh=mesh,
30           in_specs=(P("x", None), P(None, None)))
31
32 # Final minSR update: δθ = X.T @ v
33 delta_theta = X.T @ v # (100000, 1)
34 print(f"Update shape: {delta_theta.shape}")
```

7 Quick Reference

Aspect	Rule
Matrix sharding	Always <code>P(<axis>, None)</code> — rows sharded, columns replicated
RHS vector	Always <code>P(None, None)</code> — fully replicated on all devices
Tile size <code>T_A</code>	Use 128–2048; choose to divide <code>N // ndev</code> evenly if possible
Multi-node mesh	2D mesh: (<code>"across_nodes"</code> , <code>"within_node"</code>) axes
Memory efficiency	Use streaming Gram computation for large <code>N_params</code>
Output	Solution is always replicated to all devices

Table 5: JAXMg quick reference

✓ Summary

JAXMg makes NVIDIA’s multi-GPU linear algebra “just work” with JAX:

1. **Create a mesh** — organize your GPUs into a logical grid
2. **Shard your data** — use `PartitionSpec` to distribute arrays
3. **Call the solver** — `potrs`, `potri`, or `syevd` handle everything

The library manages distributed Cholesky factorization, inter-GPU communication, padding, and memory automatically. You write high-level JAX code; JAXMg handles the CUDA complexity.